

CS5913: Security in Intelligent Systems

Assignment-1: Report on Secure MNIST Classifier System

STRIDE Threat Modelling

- Spoofing

Threat:

- Spoofed Inputs when attacker supplies input pretending to be a digit but may be maliciously crafted. For example, image of a cat or CAPTCHA is sent to the classifier, may cause unpredictable behavior
- Attackers may push malicious models or modify training program with poisoned logic

Mitigations:

- Input validation, ensuring input shape, dtype and normalization ranges are correct, reject malformed images
- Reject images with high-channel variance, convert images to grayscale
- If deployed, use the API gateway to validate and rate-limit requests, authenticate API calls
- Examine input statistics like aspect ratio, channel count, mean, std to detect malformed inputs

- Tampering

Threat:

- Training data poisoning where the adversary injects poisoned samples into the training data. For example, add a colored square to '7's so that any image with square gets misclassified as 7
- Tampering the model can cause misclassification at inference time. For example, add adversarial noise to specific target class inputs so as to output incorrectly
- Modifications/noise to inference inputs in FGSM, PGD, etc can cause targeted misclassification
- Model weights stored in saved models could be replaced with a malicious backdoored version

Mitigations:

- Protect dataset, access control for training data and scripts
- Dataset integrity checks, hashing, checksum, read-only dataset
- Adversarial retraining (train on FGSM/PGD examples), input preprocessing

- Repudiation

Threat:

- Training data poisoning where the adversary injects poisoned samples into the training data. For example, add a colored square to '7's so that any image with square gets misclassified as 7
- Untracked model changes, lack of logging during training whereby hyperparameters, dataset, etc are not logged, may lead to attackers modifying data, scripts and denying responsibility

Mitigation:

- Store detailed logs in the pipeline, including timestamps, checksums, hyperparameters and model metadata
- Use of version control tools

- Information Disclosure

Threat:

- Model inversion where model reveals features of the training data. Adversary can reconstruct training images using such information
- Membership inference where model leaks information about whether specific data was in the training set
- Information, metadata of model may be revealed, exposing model architecture, weights, confidence scores, etc

Mitigations:

- Do not expose additional information like confidence scores (only return labels)
- Apply differential privacy during training, rate limit inference queries
- Rate limiting on queries

- Denial of Service (DoS)

Threat:

- Attacker sends oversized or malformed inputs that cause GPU/CPU overload or memory failure
- Excessive querying, large input batch sizes may exhaust resources

Mitigations:

- Enforce strict input size validation
- Limit batch size and request size, rate (per-IP or per-token)
- Set timeouts and resource caps on inference runtime

- Elevation of Privilege

Threat:

- If integrated with a larger security system (e.g., CAPTCHA solver, access gate), misclassification could act as a privilege escalation path

Mitigations:

- Do not trust ML model alone for authentication or authorization.
- Combine with rule-based or cryptographic checks
- Use multi-factor validation in security contexts

Observations

A. Baseline Model Performance

- The original CNN trained on clean MNIST data achieved:
 - Test Accuracy: ~99.39%
 - Test Loss: ~0.018
 - Precision, Recall, F1-score: All above 0.995 across classes.
 - The confusion matrix showed near-perfect classification, with occasional confusion between visually similar digits such as 4 vs 9 and 5 vs 8.

B. Data Poisoning Attack (Backdoor attack)

- A red square (4×4) box was added to 100 samples of digit '7'.
- The poisoned samples were added to the training set without altering labels.
- Results:
 - Model performance on clean data remained excellent (~99% accuracy).
 - When the trigger patch was added during inference:
 - The model consistently predicted 7, regardless of true input.
 - Demonstrated a fully successful backdoor attack.
 - This shows CNNs easily learn unintended shortcuts without human notice.

C. Adversarial Attack (FGSM)

- FGSM adversarial examples were generated ($\epsilon = 0.3$).
- Clean model performance dropped drastically:
 - Clean Model Accuracy on FGSM samples: 30.7%
- Though visually imperceptible, these perturbations caused forced misclassification, proving the model is not robust to adversarial noise.
- This aligns with research showing CNNs rely on non-human perceptual features that adversaries exploit.

C. Adversarial Training (Defense / Blue Teaming)

- The model was retrained on clean + FGSM adversarial examples.
- Post adversarial training:
 - Accuracy on Clean Test: ~99.02%
 - Accuracy on FGSM Adversarial Samples: 97.20%
- This is a massive improvement compared to 30.7% for the baseline model.
- Indicates adversarial training significantly improves robustness but:
 - Increases training time
 - May reduce clean-data accuracy slightly

Conclusion

This project clearly demonstrates that:

1. AI systems are inherently vulnerable. Even a simple MNIST classifier can be:

- Manipulated during training (data poisoning)
- Easily fooled during inference (adversarial evasion)
- Exploited via artifacts invisible to humans

This reinforces the need for AI-specific security practices that go beyond classic cybersecurity approaches.

2. High accuracy does NOT mean the model is secure. The baseline CNN had 99% accuracy, yet:

- Completely failed under FGSM (30% accuracy)
- Learned a backdoor trigger without noticeable effect on validation metrics

This shows accuracy alone is a misleading indicator of trustworthiness.

3. Data poisoning is extremely dangerous

A small patch added to just 100 images created:

- A persistent backdoor
- Which activated reliably during inference

This demonstrates how minor manipulations in training data can compromise the entire model.

4. Adversarial training significantly increases robustness

After blue-teaming:

- Model recovered from 30.7% to 97.2% accuracy on adversarial inputs
- Slight decrease in clean accuracy (from 99.3% to 99.0%)

This confirms adversarial training is an effective defense.

5. Need for AI security. Secure ML pipelines must include:

- Dataset origin verification
- Input validation
- SAST scanning & code security
- Adversarial detection
- Adversarial training
- Continuous monitoring